

AI Guardrails Staff Starter Pack

Classify proposed administrative AI uses, set information boundaries and make accountable human review visible.

What this helps you do

Give staff a usable first screen for routine AI-assisted work while keeping tool approval, sensitive information, high-stakes decisions and escalation with the organisation's authorised owners.

Companion workbook

Open **ai-guardrails-staff-starter-pack.xlsx** to enter ratings, use formulas and keep a working copy in an approved location.

Information boundary

Use general, synthetic or appropriately de-identified examples. Do not enter participant, patient, worker, clinical, credential or other personal or sensitive information into a public download or unapproved system.

BEFORE ANY PROMPT

Use a five-step decision path

1	Tool and use Is this tool and this specific use case approved by the organisation? Tool approval alone is not use-case approval.
2	Information Can the task be completed without personal, sensitive, confidential, credential or client information? If no or unsure, stop before entering anything.
3	Decision boundary Is this a low-risk assistive task rather than a clinical, legal, employment, safety or other high-stakes decision or recommendation?
4	Human review Is a qualified human reviewer named, with clear checks and authority to edit or reject the output?
5	Control and record Is there an escalation path, approved storage location, recorded decision and review date?

When uncertain

Stop before entering information. Use an approved non-AI process or ask the organisation's authorised AI, privacy, security, clinical, legal or other qualified owner as applicable.

Screening is not approval

The workbook can identify a candidate, conditional or prohibited-pending-review path. An authorised person still records the actual decision and conditions.

STARTER RULES

Turn principles into staff rules

Rule	Staff action	Owner evidence
Approved tool and use case	Use only the tool and task that the organisation has approved.	Approval source, owner and review date.
Minimum permitted information	Enter only what is necessary and allowed for the approved purpose.	Information boundary and examples.
No high-stakes delegation	Do not let AI make clinical, legal, employment, safety or other material decisions.	Named accountable decision-maker.
Human review	Check facts, omissions, sources, uncertainty, tone and context before use.	Reviewer role and acceptance criteria.
Escalation	Stop and escalate errors, unsafe outputs, uncertainty or out-of-scope uses.	Named escalation path and response owner.
Approved storage	Store prompts, drafts and final outputs only under the organisation's rule.	Approved location, access and retention rule.
Ongoing review	Reassess when the tool, model, data, purpose, risk or process changes.	Next review date and change trigger.

Typical screening examples

- + Candidate for approval: first draft of a generic internal checklist, using an approved tool, no sensitive information and a named reviewer.
- + Conditional: generic procedure wording where the use is suitable but the tool approval, storage rule or reviewer is still unclear.
- + Prohibited pending review: sensitive information in an unapproved tool, or AI making or substituting for a high-stakes decision.

STAFF CHECKLIST

Check before you prompt

Check	Before entering anything
☐	The task and expected output are clear.
☐	The tool and this use case are approved for this purpose.
☐	Only the minimum permitted information will be entered.
☐	No personal, sensitive, confidential or credential information is included.
☐	Reliable source material is available and permitted for this use.
☐	A qualified human reviewer is named and can reject the output.
☐	The escalation path and approved storage location are known.

Prompting discipline

- + State the task, audience, permitted source material and required format.
- + Do not ask the model to invent missing facts, authority or evidence.
- + Require uncertainty to be visible and important claims to be checked against approved sources.
- + Keep the qualified person responsible for the decision and final use.

Public tools

Do not enter personal or sensitive information into a publicly available or unapproved AI tool. Your organisation must complete its own due diligence and approve the tool, use, information and controls.

HUMAN REVIEW

Review the output and close the task

Check	Before using the output
☐	Facts, names, dates, amounts and references are checked.
☐	Important omissions, contradictions and uncertainty are addressed.
☐	Claims are supported by approved source material.
☐	Tone, audience and context are appropriate.
☐	The output does not imply certainty or authority it does not have.
☐	The qualified reviewer accepts, edits or rejects the output.
☐	Any error, unsafe output or near miss is escalated as required.

Close the task

- + Record the final version and reviewer decision.
- + Store the final material only in the approved location.
- + Retain or delete prompts, drafts and working files under the approved rule.
- + Keep conditions, exceptions and the next review date visible.
- + Route lessons, incidents or near misses to the responsible owner.

Reviewer authority

Human review is meaningful only when the reviewer is qualified for the decision, has enough evidence and is allowed to reject the output.

OPERATING CONTROL

Maintain a simple use-case register

Register field	What to record
Use case and purpose	What staff want AI to assist with and why.
Tool approval	Yes, No or Unsure, with the approval source.
Information type	Whether personal, sensitive, confidential or credential information is involved.
Decision risk	Whether the output affects a high-stakes or qualified decision.
Human reviewer	Named qualified role, checks and authority to reject.
Escalation and storage	Where uncertainty goes and where permitted material is stored.
Classification	Candidate for approval, Conditional - review required, or Prohibited pending review.
Decision and review	Authorised outcome, conditions, owner, date and next review trigger.

Not enough for every use case

This pack is a starter operating template. It does not replace legal, privacy, clinical, employment, security, regulatory or other qualified review, and it does not transfer accountability away from the organisation.

SERVICE BRIDGE

Choose the next level of work

What the starter pack reveals	Possible next step	Boundary
A few low-risk administrative uses are clear, with approved tools, information boundaries and accountable review.	Record formal organisational approval and train staff on the final rules.	The starter pack does not itself approve the use.
Many uses are already happening, ownership is unclear or staff cannot apply the rules to real tasks.	Heutrix AI Guard may help create a usable register, boundaries, review, escalation and staff guidance.	It is an operational service, not legal, privacy, clinical or compliance advice.
The issue is a wider workflow problem rather than AI use.	Use the Workflow Bottleneck Scorecard or discuss Heutrix Diagnostics / Workflow Transformation as appropriate.	Start with the workflow, not with AI.
A use involves sensitive information or high-stakes decisions.	Stop and use qualified organisational review and an approved non-AI path unless explicitly authorised.	Do not use a public or unapproved tool.

Initial conversation

Talk to Heutrix: a 20-minute introductory conversation about one high-level workflow.

Keep it general. Do not include sensitive records. Any paid work is agreed separately.